


End-to-End Training and Inference Pipeline



Training: With augmentation + Label Smoothing Loss + AdamW

Inference: No augmentation, ensemble averaging (16f + 32f)